

Why Use Concurrency?

For decades, software engineers had a "**free lunch**." We could write standard, sequential code (one instruction at a time), and every two years, our code would magically run faster simply because **Intel** or **AMD** released a faster processor.

That era is officially over. To understand why we must now write concurrent code, we have to look at the physical limitations of modern hardware.

1. The Death of Moore's Law and the "Free Lunch."

Back then, **Moore's Law** was an observation that the number of transistors on a microchip **doubled** roughly **every two years**. Smaller transistors meant faster switching times (higher clock speeds), leading to exponential performance growth.

However, we have hit physical limits that prevent us from continuously shrinking transistors and increasing the clock speed **f (frequency)**.

2. The Power Wall and Dennard Scaling

Why can't we just make a **10 GHz** or **20 GHz** processor? The answer is **heat** and **power** consumption. The dynamic power consumption of a processor is defined by this equation:


$$P = \alpha * C f V^2$$

(Where **alpha** is the switching activity, **C** is capacitance, **f** is clock frequency, and **V** is voltage).

- **The Problem:** If you increase the **frequency (f)**, power **consumption (P)** goes up. High power creates physical heat. If we push the clock speed too high, the silicon chip will literally **melt**.
- **Dennard Scaling:** For years, engineers kept power low by shrinking the **voltage (V)** as transistors got smaller. Because voltage is squared (**V²**), dropping the voltage saved massive amounts of power.
- **The Wall:** We can no longer drop the voltage. If the voltage gets too low, the processor cannot distinguish between a real "**1**" or "**0**" signal and random electrical background noise. Additionally, as insulators get microscopic, **Leakage Power** (electricity bleeding through the chip even when it isn't switching) creates constant heat.

3. The Von Neumann Bottleneck

Even if we increase the CPU clock speed without melting the chip, we would hit a data traffic jam.



In the **Von Neumann architecture**, the **CPU** must constantly fetch instructions and data from Main Memory (**RAM**). Main Memory is slow compared to a modern **CPU**. Increasing the **CPU** speed doesn't help if the **CPU** spends **80%** of its time sitting idle, waiting for the **RAM** to deliver the next piece of data.

4. The Solution: Multi-Core Processors

Because hardware engineers can no longer increase the raw clock speed of a single processor, they changed their strategy: **Replicated Hardware**.

Instead of giving you one insanely fast processor, they now give you **4, 8, or 16** slightly slower processors (Cores) glued together on a single chip.

5. Parallel Execution vs. Sequential Tasks

To actually use a multi-core processor, your code must execute in **Parallel** (multiple instructions running at the exact same physical instant on different cores).

- **The Catch:** Not all tasks can be parallelized.
- **The Dishwashing Analogy:** If you have one sink and one drying rack, two people cannot work at the same time. Furthermore, you cannot dry a dish before it is washed. Some logic is strictly sequential.
- **The Programmer's Burden:** Automated compilers are terrible at figuring out how to safely split a sequential program across multiple cores. It is now the software engineer's job to explicitly write concurrent code to divide the tasks. If you write standard, sequential code on an **8-core** machine, your program will only run on **1 core**, and the other **7** will sit completely idle.

Expert Takeaway

This hardware shift is the exact reason Google created **Go**. Languages like **C**, **C++**, and **Python** were designed in the **single-core era**, meaning concurrent programming was tacked on later using clunky, difficult-to-manage external libraries. Go was designed from day one for the **multi-core** era. By baking concurrency directly into the syntax of the language, Go makes it incredibly easy to split tasks across all available CPU cores, squeezing every ounce of performance out of modern hardware!